

Highlights

Do Explanation Signatures Detect Classification Errors? A Controlled Comparison of Graph and Token Explainers

Fabio Yáñez-Romero, Andrés Montoyo, Armando Suárez, Yoan Gutiérrez, Ruslan Mitkov

- A protocol to compare graph and token explainers on the same model and budget.
- Explanation features add nothing to a language model's confidence as error signals.
- Within the budget every explanation is sufficient and not necessary, graph or token.
- Distilled surrogates share errors with gold-trained twins but are confident in them.
- Explanations help a distilled surrogate only where its confidence was damaged.